

Summary: "How to Read a Paper" by S. Keshav

Keshav points out something pretty relatable: researchers spend hundreds of hours a year reading papers, but nobody ever actually teaches you how to do it well. Most people just figure it out through trial and error, which wastes a lot of time and gets frustrating fast. His solution is simple read a paper in **three passes**, each one going deeper than the last, instead of reading start to finish in one go.

Pass 1 : Just read the title, abstract, intro, headings, and conclusion. Skim the references. By the end, you should be able to answer the "five Cs" what Category the paper is, its Context (related work), whether the assumptions seem Correct, what its main Contributions are, and whether it's written with Clarity. This tells you if it's even worth reading further.

Pass 2 : Read more carefully, look closely at figures and graphs (check if axes are labeled properly, results have error bars, etc.), and jot down key points. By the end, you should be able to explain the paper's main idea and evidence to someone else, even if you don't follow every detail.

Pass 3 : Try to mentally "re-implement" the paper recreate the work using the same assumptions the authors made, then compare your version to theirs. This is where you really catch hidden flaws, unstated assumptions, or things you'd have done differently.

He also explains how to use this method for literature surveys: skim a handful of recent papers, mine their reference lists for repeated citations to find the key papers/researchers in the field, then check where those researchers publish to find the top venues.

What I like most is the permission to stop early not every paper deserves a full read, and deciding that after just 10 minutes saves a ton of time. Next time I read a paper, I plan to actually time-box pass one and force myself to answer the five Cs before deciding whether to go deeper.